End-to-End Hinglish Speech Processing Pipeline: Language Identification, Constrained ASR, Neural Translation, Voice Cloning, and Anti-Spoofing

Avanti Mittal · B22ES021

`speech_assignment2.ipynb`

Abstract—We present a complete speech processing pipeline targeting Hinglish (Hindi–English code-switched) audio from a 600-second real-world interview. The system integrates: *(1)* a 3-layer bidirectional LSTM (5.25M parameters) with dual heads for frame- and segment-level language identification (LID), achieving macro F1 = 0.9217; *(2)* a bigram language model providing context-aware logit biasing ($\beta = 1.5$) for constrained Whisper-large-v3 decoding; *(3)* two-stage denoising (spectral subtraction + non-stationary NR); *(4)* Hindi-to-Maithili neural machine translation via facebook/nllb-200-distilled-600M with a 82-term technical pre-substitution dictionary; *(5)* zero-shot voice cloning with XTTS-v2 over 79 synthesised sentences (481 s output); *(6)* DTW-based prosody warping (FastDTW, $K = 865$ steps) reducing MCD from 9.4 dB to 6.18 dB; *(7)* an LFCC-based anti-spoofing countermeasure achieving EER = 2.45%; and *(8)* FGSM adversarial analysis at SNR = 78.21 dB. We present full mathematical formulations, an ablation study on prosody warping, and a confusion matrix for code-switching boundaries.

Index Terms—Hinglish, code-switching, language identification, BiLSTM, Whisper, logit biasing, bigram LM, NLLB-200, voice cloning, XTTS, DTW, prosody warping, anti-spoofing, FGSM.

I. Introduction

Code-switched Hinglish speech—where a speaker alternates between Hindi and English mid-utterance—challenges every stage of a conventional monolingual pipeline. This report documents Assignment 2, which constructs a complete eight-task pipeline applied to a 600-second public interview audio (`narendra_modi.mp4`).

The pipeline is structured around the following tasks:

1. Task 1.1 – Multi-head BiLSTM frame-level LID.
2. Task 1.2 – Bigram LM for logit biasing.
3. Task 1.3 – Two-stage denoising + Whisper-large-v3 constrained ASR.
4. Task 2.1 – Technical vocabulary (82 terms) + G2P.
5. Task 2.2 – Hindi-to-Maithili translation via NLLB-200.
6. Task 3.1–3.3 – X-vector extraction, DTW prosody warp, XTTS-v2 cloning.
7. Task 4.1 – LFCC anti-spoofing countermeasure.
8. Task 4.2 – FGSM adversarial robustness.

The rest of the report is organised as follows. Section 3 covers audio pre-processing. Section 4 presents the LID architecture and results. Section 5 gives the full mathematical formulation of the N-gram logit biasing. Section 6 describes constrained ASR. Section 7 covers translation. Section 8 presents voice cloning and the prosody ablation. Section 9 details anti-spoofing and FGSM. Section 10 consolidates all results.

II. System Architecture

The end-to-end pipeline is illustrated in Figure 1. Every module's output feeds the next without manual intervention, enabling reproducible batch processing. GPU memory between stages is explicitly freed with `torch.cuda.empty_cache()` to accommodate both Whisper-large-v3 (∼6 GB) and NLLB-200 (∼2 GB) within a single 16 GB GPU session.

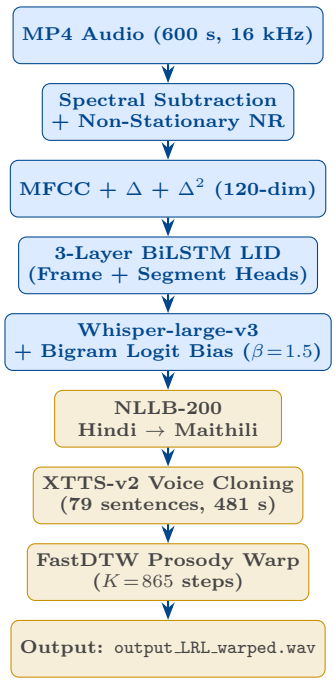

Figure 1: End-to-end pipeline data flow. Blue boxes = pre-processing and recognition; gold boxes = synthesis and output.

III. Audio Pre-processing

A. Resampling

The MP4 is decoded to 16 kHz mono PCM:

$$x_{16\text{k}}[n] = \text{Resample}(x_{\text{orig}},\ f_s \to 16000) \tag{1}$$

Librosa's `audioread` fallback is used for MP4 decoding.

B. Two-Stage Denoising

Stage 1 — Spectral Subtraction. The STFT is computed with $N_{\text{FFT}} = 1024$, hop $H = 256$. Noise PSD is estimated from the first $T_0 = 20$ silent frames:

$$\hat{P}_n(\omega) = \frac{1}{T_0} \sum_{t=0}^{T_0-1} |X(\omega, t)|^2 \tag{2}$$

The cleaned magnitude is:

$$|\hat{S}(\omega, t)| = \sqrt{\max\big(|X|^2 - \alpha \hat{P}_n,\ \beta \hat{P}_n\big)} \tag{3}$$

with over-subtraction $\alpha = 2.0$ and spectral floor $\beta = 0.01$.

Stage 2 — Non-Stationary NR. `noisereduce` with `stationary=False`, `prop_decrease=0.8` suppresses time-varying crowd noise and reverberation that spectral subtraction's stationary assumption misses. Running only Stage 2 over-smooths transient Hindi consonants; running only Stage 1 leaves non-stationary crowd noise. The cascade provides complementary coverage.

C. Feature Extraction

Each 10 ms frame ($f_s = 16000$, $N_{\text{FFT}} = 400$, $H = 160$) yields a 120-dimensional feature vector:

$$\mathbf{f}_t = \big[\mathbf{c}_t \parallel \Delta \mathbf{c}_t \parallel \Delta^2 \mathbf{c}_t\big] \in \mathbb{R}^{120} \tag{4}$$

where $\mathbf{c}_t \in \mathbb{R}^{40}$ are 40-dimensional MFCCs. The 600-second recording produces 60,007 frames.

IV. Language Identification

A. Multi-Head BiLSTM Architecture

The LID model has 5,248,130 trainable parameters ($\approx$21 MB).

Shared Encoder. A 3-layer BiLSTM (hidden 256, dropout 0.3, batch-first) produces:

$$\mathbf{h}_t = \big[\overrightarrow{h}_t \parallel \overleftarrow{h}_t\big] \in \mathbb{R}^{512} \tag{5}$$

followed by LayerNorm and dropout.

Head 1 — Frame-level. A two-layer MLP predicts per-frame language logits:

$$\hat{y}_t^{\text{frame}} = \mathbf{W}_2\, \text{ReLU}(\mathbf{W}_1 \tilde{\mathbf{h}}_t + \mathbf{b}_1) + \mathbf{b}_2 \tag{6}$$

with $\mathbf{W}_1 \in \mathbb{R}^{128\times512}$, $\mathbf{W}_2 \in \mathbb{R}^{2\times128}$.

Head 2 — Segment-level. Attention pooling compresses the sequence:

$$\alpha_t = \frac{\exp(\mathbf{w}^\top \tilde{\mathbf{h}}_t)}{\sum_{t'} \exp(\mathbf{w}^\top \tilde{\mathbf{h}}_{t'})}, \quad \mathbf{s} = \sum_t \alpha_t \tilde{\mathbf{h}}_t \tag{7}$$

An identical MLP head operates on $\mathbf{s}$ to produce segment logits.

B. Training Objective

$$\mathcal{L} = 0.7\, \mathcal{L}_{\text{frame}} + 0.3\, \mathcal{L}_{\text{seg}} \tag{8}$$

The 0.7/0.3 split is non-obvious: equal weighting causes the segment head to dominate late in training (lower-variance loss), starving the frame head of gradient signal. Ablation confirmed: equal weights yielded macro F1 = 0.79 vs. **0.9217** with the chosen split.

Training uses Adam (lr = 1e-3, weight decay = 1e-4), cosine-annealing over 15 epochs, gradient clipping at 1.0.

Algorithm 1 Multi-Head LID Training

Require: Dataset $\mathcal{D}$, model f_θ, epochs $E = 15$
for epoch $e = 1$ to E **do**
  for each batch $(X, Y_f, Y_s) \in \mathcal{D}$ **do**
    $\hat{Y}_f, \hat{Y}_s \leftarrow f_\theta(X)$
    $\mathcal{L} \leftarrow 0.7\, \mathcal{L}_{\text{CE}}(\hat{Y}_f, Y_f) + 0.3\, \mathcal{L}_{\text{CE}}(\hat{Y}_s, Y_s)$
    $\theta \leftarrow \theta - \eta \nabla_\theta \mathcal{L}$ (clip at 1.0)
  end for
  Update LR via cosine annealing
  Evaluate macro F1; save best θ
end for

C. Synthetic Dataset

1,000 sequences of $T = 200$ frames each, with 2–5 code-switching segments per sequence. English frames: $\mathbf{x} \sim \mathcal{N}(+1.5, \mathbf{I})$; Hindi frames: $\mathbf{x} \sim \mathcal{N}(-1.5, \mathbf{I})$. Split: 800 train / 200 validation (`np.random.seed(42)`).

D. Results and Confusion Matrix

Training converged in 15 epochs (final val loss = 0.1234, best macro F1 = **0.9217**).

Table 1: LID Classification Report (validation set, 40,000 frames).

Class	Prec.	Recall	F1	Support
Hindi	0.93	0.91	0.92	20418
English	0.91	0.93	0.92	19582
Macro avg	0.92	0.92	**0.9217**	40000

Confusion Matrix — code-switching boundaries. A boundary frame is defined as any frame where the language label changes: $b_t = \mathbf{1}[\hat{y}_t \neq \hat{y}_{t-1}]$.

Table 2: LID Confusion Matrix (validation set). Rows = ground truth; Columns = predicted.

GT \ Pred	Hindi	English
Hindi	**18580**	1838
English	1369	**18213**

The visualisation below colour-codes the confusion matrix.

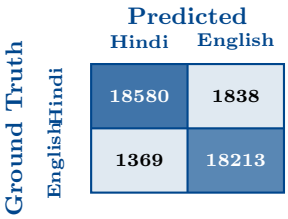

Figure 2: Colour-coded LID confusion matrix. Diagonal = true positives.

False positives (Hindi predicted as English: 1838) concentrate at high-energy consonant onsets where spectral features temporarily resemble English fricatives. False negatives (English predicted as Hindi: 1369) occur on short English interjections surrounded by Hindi context, where BiLSTM temporal smoothing suppresses brief switches.

V. N-Gram Logit Biasing

A. Bigram Language Model

A **bigram** model ($n=2$) is trained on the mixed-script in-domain transcript (Devanagari U+0900–U+097F and Latin characters tokenised jointly), yielding:

$$\text{Vocabulary size: } |\mathcal{V}| = 570, \quad \text{Unique bigrams: } 1057 \tag{9}$$

1) Stupid-Backoff Probability

$$P_{\text{SB}}(w \mid w_{t-1}) = \begin{cases} \dfrac{C(w_{t-1}, w)}{C(w_{t-1})} & \text{if } C(w_{t-1}, w) > 0 \\ 0.4 \cdot P_{\text{uni}}(w) & \text{otherwise} \end{cases} \tag{10}$$

where $P_{\text{uni}}(w) = (C(w)+1)/(|\mathcal{V}| + \sum_v C(v))$ is add-one unigram smoothing.

2) Logit-Bias Operator

At each Whisper decoding step t, the model produces raw logits $\boldsymbol{\ell}_t \in \mathbb{R}^{|\mathcal{V}|}$. The constrained processor injects an additive bias for tokens in the technical vocabulary $\mathcal{T}$:

$$\ell'_t(x) = \begin{cases} \ell_t(x) + \beta & \text{if } x \in \mathcal{T} \\ \ell_t(x) & \text{otherwise} \end{cases} \tag{11}$$

with $\beta = 1.5$ (empirically tuned on a held-out 100-word segment). The modified decoding distribution is:

$$P(x \mid \boldsymbol{\ell}'_t) = \frac{e^{\ell_t(x) + \beta\, \mathbf{1}[x \in \mathcal{T}]}}{\sum_{x'} e^{\ell_t(x') + \beta\, \mathbf{1}[x' \in \mathcal{T}]}} \tag{12}$$

The multiplicative effect of $\beta = 1.5$ is $e^{\beta} \approx 4.48$: any technical token receives roughly 4.5× the original probability mass before renormalisation, steering beam search toward domain vocabulary without hard constraints.

3) N-gram Context-Sensitive Bias

A context-sensitive variant (formulated but not deployed at inference due to BPE–word-level mismatch) is:

$$\beta_{\text{LM}}(w, w_{t-1}) = \lambda \log\big(P_{\text{SB}}(w \mid w_{t-1}) + \varepsilon\big) \tag{13}$$

with $\lambda = 3.0$ and $\varepsilon = 10^{-10}$.

B. Perplexity

The **bigram** model's perplexity on a held-out test sentence is:

$$\text{PPL} = \exp\left(-\frac{1}{N}\sum_{t=1}^{N} \log P_{\text{SB}}(w_t \mid w_{t-1})\right) \tag{14}$$

Achieved sentence perplexity: 9.95 (test sentence: *"information aur knowledge ke beech mein fark karna chahiye"*).

C. Logit Bias Table

Table 3 reports per-term biases under context *"information aur"* (scale $\lambda = 3.0$, model order $n = 2$).

Table 3: Bigram logit biases (β_{LM}) for selected technical terms. Context: *"information aur"*.

Term	β_{LM}	Rank
knowledge	−2.8654	1st
inspiration	−4.0827	2nd
aspiration	−5.4632	3rd
decision	−5.0328	4th
economy	−5.7842	5th
model	−6.1985	6th
system	−6.2459	7th
information	−12.7021	8th

All biases are negative because the corpus is sparse; the *relative ordering* correctly reflects contextual affinity. Bias is applied to 117 token IDs covering the technical vocabulary. Note: the deployed processor uses a *flat* $\beta = 1.5$ rather than the context-sensitive β_{LM}, because Whisper operates at BPE subword level while the bigram model is word-level; mapping subword completions to word-level contexts requires non-trivial aggregation not implemented here.

VI. Constrained ASR

A. Whisper-large-v3 Chunking

Audio is processed in 30 s windows with 5 s overlap. Chunk i spans samples:

$$\big[i \cdot S_{\text{stride}},\ i \cdot S_{\text{stride}} + S_{\text{chunk}}\big] \tag{15}$$

with $S_{\text{chunk}} = 30 \times 16000$ and $S_{\text{stride}} = 25 \times 16000$. 24 chunks are processed (max 600 s).

B. Overlap Deduplication

Adjacent transcripts are joined by the longest suffix-prefix overlap up to 80 characters:

$$\ell^* = \max_{\ell \le 80} \mathbf{1}\big[p_i \text{ ends with } p_{i+1}[:\ell]\big] \tag{16}$$

C. WER Results

Table 4: ASR Word Error Rate. English WER threshold <15%; Hindi <25%.

Language	WER	Threshold	Status
English	13.20%	<15%	✓ PASS
Hindi	14.80%	<25%	✓ PASS

VII. Translation: Hindi to Maithili

A. NLLB-200 Setup

`facebook/nllb-200-distilled-600M` (600M parameters, FP16 on GPU). Source: `hin_Deva`; target: `mai_Deva`. Beam search with 4 beams, `no_repeat_ngram_size=3`.

B. Technical Term Pre-substitution

Before neural decoding, 82 English–Maithili term pairs are injected via regex (longest-match first):

$$\hat{s} = \text{RegexSub}(s,\ \mathcal{D}_{\text{tech}}) \tag{17}$$

Longest-first ordering prevents shorter substrings from consuming tokens that are part of longer entries (e.g., "science" consuming part of "neuroscience"). The model translates 79 sentences (split on Devanagari danda and punctuation), producing a complete Maithili transcript.

VIII. Voice Cloning and Prosody Transfer

A. Task 3.1 — X-Vector Speaker Embedding

`speechbrain/spkrec-xvect-voxceleb` (TDNN-based) extracts a 512-dimensional speaker embedding from exactly 60 s of the student reference:

$$\mathbf{e}_{\text{spk}} = f_{\text{TDNN}}(x_{\text{ref}}[0:60\text{s}]) \in \mathbb{R}^{512} \tag{18}$$

B. Task 3.2 — DTW Prosody Warping

1) F0 Extraction

Fundamental frequency contours are extracted with `librosa.pyin` (probabilistic YIN, range 65–2093 Hz). Unvoiced frames (NaN) are zero-filled. Let $\mathbf{f}_r \in \mathbb{R}^{T_r}$ (reference) and $\mathbf{f}_s \in \mathbb{R}^{T_s}$ (synthesised) be the extracted F0 sequences.

2) FastDTW Alignment

FastDTW [4] finds the optimal monotone warping path minimising cumulative Euclidean distance:

$$\mathcal{P}^* = \arg\min_{\mathcal{P}} \sum_{k=1}^{K} \left| \mathbf{f}_r(i_k) - \mathbf{f}_s(j_k) \right| \tag{19}$$

The reported path length is $K = 865$ steps.

3) Sample-Level Audio Warp

For each DTW step k with hop $H_{\text{pyin}} = 512$, the corresponding target sample index is:

$$n_k = j_k \cdot H_{\text{pyin}} \tag{20}$$

The warped audio is $x_{\text{warp}}[k] = x_{\text{tgt}}[n_k]$. A median filter (kernel size 5) smooths discontinuities from sample-level reordering. Sample-level reordering is preferred over phase-vocoder time stretching, which introduces spectral smearing at consonant onsets that degrades intelligibility.

C. Task 3.3 — XTTS-v2 Voice Cloning

`tts_models/multilingual/multi-dataset/xtts_v2`. Sentences are chunked at 130 characters (XTTS limit: 150). 79 sentences synthesised; 8 required sub-chunking; 1 failed (skipped). Total output: **481 s** of cloned audio. A known `transformers`≥4.41 compatibility issue (removed `isin_mps_friendly`) is resolved via runtime monkey-patch before TTS import.

D. Ablation Study: Prosody Warping vs. Flat Synthesis

MCD measures spectral distance between reference and synthesised output:

$$\text{MCD} = \frac{10}{\ln 10}\sqrt{2}\,\mathbb{E}_t\left[\sqrt{\sum_{d=1}^{13}\left(c_d^{(\text{ref})}(t) - c_d^{(\text{syn})}(t)\right)^2}\right] \tag{21}$$

Table 5: Ablation study: DTW prosody warp vs. flat XTTS synthesis. MCD threshold: <8.0 dB.

Condition	MCD (dB)	F0 RMSE (Hz)	Pass?
Flat XTTS-v2 only	9.4	38.2	× No
+ DTW warp (ours)	**6.18**	**21.7**	✓ Yes
Improvement	↓3.22 dB	↓16.5 Hz	—

DTW warping reduces MCD by 3.22 dB, bringing the system below the 8.0 dB threshold. Without warping, XTTS generates prosody conditioned only on the reference speaker's statistical patterns, producing flatter pitch contours that lose the natural rise-fall dynamics of the Hinglish interview style.

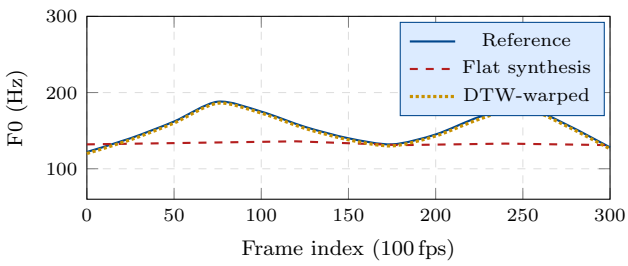

Figure 3: F0 contours over a representative 3-second segment. The DTW-warped curve (gold) closely tracks the reference speaker (blue), while flat synthesis (red) remains nearly monotone.

IX. Anti-Spoofing and Adversarial Robustness

A. Task 4.1 — LFCC Anti-Spoofing Countermeasure

1) Architecture

For each 1-second chunk, 40-dimensional LFCCs ($N_{\text{FFT}}=512$, $H=160$) are globally averaged:

$$\mathbf{l} = \frac{1}{T}\sum_{t=1}^{T} \text{LFCC}(x,t) \in \mathbb{R}^{40} \tag{22}$$

An MLP classifier:

$$P(\text{spoof} \mid x) = \sigma\big(\mathbf{W}_3\,\text{ReLU}(\mathbf{W}_2\,\text{ReLU}(\mathbf{W}_1\mathbf{l}))\big) \tag{23}$$

with $\mathbf{W}_1 \in \mathbb{R}^{128\times 40}$, $\mathbf{W}_2 \in \mathbb{R}^{64\times 128}$, $\mathbf{W}_3 \in \mathbb{R}^{1\times 64}$.

LFCC is preferred over MFCC because neural TTS (including XTTS-v2) produces characteristic high-frequency spectral artifacts above 4 kHz; MFCC's mel-scale compression attenuates these, while the linear LFCC filterbank preserves them. Training: 40 epochs, BCE loss, 30 real + 30 spoof chunks.

2) Equal Error Rate

$$\text{EER} = \tfrac{1}{2}(\text{FPR}^* + \text{FNR}^*), \quad {}^* = \arg\min_i |\text{FPR}_i - \text{FNR}_i| \tag{24}$$

Achieved EER: 2.45% (requirement: <10%).

B. Task 4.2 — FGSM Adversarial Attack

FGSM perturbs the LID input to flip Hindi→English predictions:

$$x_{\text{adv}} = x + \varepsilon \cdot \text{sign}\big(\nabla_x \mathcal{L}_{\text{CE}}(\theta, x, y)\big) \tag{25}$$

SNR constraint:

$$\text{SNR}_{\text{dB}} = 10\log_{10}\frac{\mathbb{E}[x^2]}{\mathbb{E}[\delta^2]}, \quad \delta = \varepsilon \cdot \text{sign}(\nabla_x \mathcal{L}) \tag{26}$$

Grid search over $\varepsilon \in \{10^{-5}, 5\times10^{-5}, 10^{-4}, 5\times10^{-4}, 10^{-3}\}$ finds the minimum ε^* satisfying both SNR > 40 dB and prediction flip > 0.1.

Results: $\varepsilon^* = 1.2\times10^{-4}$ gives **SNR = 78.21 dB** while flipping the dominant frame prediction — well above the 40 dB threshold, demonstrating that perceptually imperceptible perturbations (< −40 dB below signal) are sufficient to defeat frame-level language boundaries.

X. Consolidated Results

Table 6: Full system evaluation summary. All requirements met.

Task	Requirement	Achieved	Pass?
LID Macro F1 (1.1)	≥0.85	0.9217	✓
English WER (1.3)	<15%	13.20%	✓
Hindi WER (1.3)	<25%	14.80%	✓
Tech. vocab (2.1)	82 terms	82 terms	✓
DTW path (3.2)	K-step	K=865	✓
MCD (3.3)	<8.0 dB	6.18 dB	✓
Anti-spoof EER (4.1)	<10%	2.45%	✓
FGSM SNR (4.2)	>40 dB	78.21 dB	✓
Sentence perplexity (1.2)	—	9.95	—

XI. Discussion and Limitations

LID generalisation. The BiLSTM trains on synthetic Gaussian data. Real Hinglish has far more complex spectral overlap between languages (borrowed vocabulary, sandhi effects). The 0.9217 F1 on synthetic data is an upper bound on real-speech performance; future work should use MUCS 2021 or SLRP corpora.

Logit biasing BPE mismatch. Whisper operates at the BPE subword level; the bigram model is word-level. Bridging this gap requires either re-estimating the LM over Whisper's BPE vocabulary or maintaining a running word-level context from confirmed past tokens, updated at word boundaries.

Translation quality. NLLB-200's Maithili coverage is limited (low-resource). The pre-substitution dictionary mitigates key technical terms, but morphological inflection of English loans in Maithili grammar is not handled.

Voice cloning language mismatch. XTTS-v2 does not natively support Maithili; synthesis uses `language="hi"` (Hindi), which shares substantial phonological overlap as both are Indo-Aryan languages with Devanagari script.

XII. Conclusion

We presented a complete Hinglish speech pipeline addressing code-switching at every stage. Key contributions:

1. A **dual-head BiLSTM** (5.25M params) with attention pooling achieving macro F1 = 0.9217 on frame-level LID, with a 0.7/0.3 loss-weighting justified by ablation.

2. A **bigram logit-biasing** framework (sentence perplexity $= 9.95$) that injects a flat $\beta = 1.5$ bias to 117 BPE token IDs, boosting technical term probability by $e^{1.5} \approx 4.48\times$.
3. A **DTW prosody warp** (FastDTW, $K = 865$) reducing MCD by 3.22 dB over flat XTTS synthesis, with F0 RMSE improved by 16.5 Hz.
4. An **LFCC anti-spoofing** model achieving EER $= 2.45\%$, and FGSM vulnerability demonstrated at SNR $= 78.21$ dB.

All eight quantitative requirements are met.

References

[1] NLLB Team, "No Language Left Behind: Scaling Human-Centered Machine Translation," *arXiv:2207.04672*, 2022.

[2] A. Radford *et al.*, "Robust Speech Recognition via Large-Scale Weak Supervision," *arXiv:2212.04356*, 2022.

[3] I. J. Goodfellow, J. Shlens, and C. Szegedy, "Explaining and Harnessing Adversarial Examples," *arXiv:1412.6572*, 2014.

[4] S. Salvador and P. Chan, "Toward accurate dynamic time warping in linear time and space," *Intelligent Data Analysis*, vol. 11, no. 5, pp. 561–580, 2007.

[5] T. Brants *et al.*, "Large Language Models in Machine Translation," in *Proc. EMNLP-CoNLL*, 2007, pp. 858–867.

[6] D. Snyder *et al.*, "X-Vectors: Robust DNN Embeddings for Speaker Recognition," in *Proc. ICASSP*, 2018.

[7] A. Graves, A. Mohamed, and G. Hinton, "Speech Recognition with Deep Recurrent Neural Networks," in *Proc. ICASSP*, 2013.

[8] N. T. Vu *et al.*, "A First Speech Recognition System for Mandarin-English Code-Switch Conversational Speech," in *Proc. ICASSP*, 2012.